



CGNS Proposal for Extension 0045

Storing Cell-Wise Polynomial Data and Generalising
the Description of Curved Grid Elements

Version 3.0

August 1, 2026

CPEX#	0045
Scope	High-order polynomial interpolation for elements and solutions
Champion	Koen Hillewaert
Organization	Cenaero
E-mail	koen.hillewaert@cenaero.be
v3 Revision Author	M. Scot Breitenfeld, The HDF Group — brtnfd@hdfgroup.org
Date First Posted	Mar. 26, 2019
Date of Revision	May 4, 2026
SIDS Status	v2 accepted; v3 amendment pending Steering Committee vote
Filemap Status	v2 accepted; v3 amendment pending Steering Committee vote
MLL Status	v2 accepted; v3 amendment pending Steering Committee vote
Reference impl.	branch CPEX45_high_order_wip

Editorial convention. Passages set in the callout shown below reproduce text *verbatim* or *near-verbatim* from CPEX-0045 v2 (K. Hillewaert, Cenaero). Longer technical passages—function-space definitions, pseudo-code listings, and key equations—are marked even when the wording is only lightly adapted from v2. Shorter paraphrased phrases drawn from v2 that are woven into v3 commentary are not individually marked. The surrounding prose (unmarked) is v3 content that formalises, extends, or clarifies the v2 material, and includes the implementation specification adopted in CGNS v5.0.

Sample: text in this style is verbatim or near-verbatim from CPEX-0045 v2.

Amendment to Approved Proposal

CPEX-0045 v2 was previously approved by the CGNS Steering Committee. This document (v3) is submitted as a revision to that approved proposal. The committee is asked to evaluate only the *delta* described below. Passages reproduced from v2 are marked throughout with a blue left-rule callout.

What v2 Established

- Two new SIDS node types: `ElementInterpolation_t` and `SolutionInterpolation_t`, both children of `Family_t`.
- `InterpolationType_t` enum with four values: `ParametricLagrange`, `ParametricMonomialsPascal`, `CartesianMonomialsPascal`, `IsoParametric`.
- Support for purely spatial and space-time interpolation via `SpatialOrder` and `TemporalOrder` fields.
- High-order solution data stored in existing `FlowSolution_t` nodes; polynomial orders recorded per solution block.
- Scope restricted to unstructured mesh computations.

Changes from v2 to v3

Each item is labelled **[NEW]** (genuinely new; requires committee vote) or **[CLARIFICATION]** (formalises v2 intent; no semantic change to the approved standard).

SIDS changes:

1. **[NEW] `LagrangeControlPointDistribution_t` enum.** Records the parametric-space distribution of Lagrange control points. Values: `GaussLobattoLegendre`, `Equidistant`, `GaussLegendre`, `WarpAndBlend`. Stored as an optional named child of `ElementInterpolation_t` or `SolutionInterpolation_t`; recommended, but not required, for `ParametricLagrange`. *Rationale:* the attribute is redundant for *interpretation* — the nodal basis is already determined by the stored control-point coordinates together with the element type and order — but naming the distribution lets a reader use a well-conditioned closed-form construction instead of inverting the Vandermonde matrix, regenerate the nodes at full working precision rather than inheriting the writer's R8 rounding, and lets `cgnscheck` cross-validate the stored coordinates.
2. **[NEW] New `GridLocation_t` enumerator `InterpolationPoints` (value 9), required for all high-order `FlowSolution_t` nodes.** This is a wire-format addition, not a reinterpretation of existing v2 semantics: v2 specified `CellCenter` as the location for high-order solutions (Section 3.4.5). v3 instead designates the new `InterpolationPoints` value as required `GridLocation` for *all* high-order `FlowSolution_t` nodes, both whole-zone (uniform-order) and subset (variable-order via `PointRange/PointList`), and for the modal types as well as Lagrange — despite its name, the enumerator labels the per-element block of degrees of freedom, not geometric points (Section 3.1.3). *Rationale:* `CellCenter` mandates a field-array length equal to the number of cells, which contradicts the high-order requirement $\sum_e N_{\text{DOFs}}(e)$ and would fail standard SIDS sizing rules in `cgnscheck`. When the parent `FlowSolution_t` carries a `PointRange` or `PointList` under `InterpolationPoints`, the listed indices are interpreted as *element indices*, and field arrays have length $\sum_e N_{\text{DOFs}}(e)$ summed over the listed elements. For backward compatibility, readers still accept legacy files that use `CellCenter` with an explicit

- PointRange/PointList (Section 6), but conformant *writers* must use `InterpolationPoints` exclusively.
3. **[CLARIFICATION] InterpolationOrders encoding.** Formalises the on-disk representation of `SpatialOrder` and `TemporalOrder` as a 2-element `IndexArray_t` child of `FlowSolution_t`. Information content is identical to v2.
 4. **[NEW] Mandatory coordinate normalisation for Cartesian modal, with normative CharacteristicLength encoding.** Requires writers to non-dimensionalise the element-local coordinates before populating `MonomialCoefficients`, and to record the factors they used in a normatively-named `CharacteristicLength DataArray_t` child of `FlowSolution_t` (Section 5) so readers can recover physical values. Two normalisations are supported and are distinguished by array rank: *isotropic* (rank 1, one factor per element) and *per-axis* (rank 2, `PhysDim` factors per element). Readers use the recorded factors and do not recompute them. *Rationale:* high-order monomials in un-normalised physical coordinates overflow at large or small element sizes, and an isotropic factor alone leaves the modal mass matrix ill conditioned on the high-aspect-ratio cells typical of wall-resolved boundary layers; per-axis scaling is the directional non-dimensionalisation established in the Taylor-basis discontinuous-Galerkin literature [6, 7]. Recording rather than mandating a formula lets a writer dump its native coefficient blocks without rescaling, per the goals of Section 1, and removes any dependence on writer and reader agreeing on a geometric computation. Both encodings and the dedicated MLL helpers `cg_sol_characteristic_length_{read,write}` ship in the v5.0 library on the `CPEX45_high_order_wip` branch.
 5. **[CLARIFICATION] CharacteristicLength is metadata, not a solution field.** States that a reader enumerating the `DataArray_t` children of a `FlowSolution_t` to build its field list must exclude `CharacteristicLength` by name, and must not apply the field-array length rule to it (Section 6.2). *Rationale:* the node is a `DataArray_t` child of `FlowSolution_t` and so collides with the SIDS convention that such children are fields. A reader that does not exclude it reports a spurious field *and* rejects conformant files, because the metadata length does not equal $\sum_e N_{\text{DOFs}}(e)$. This was a real defect in the reference implementation, invisible on test files with no element sections because field size checking cannot run in that case.
 6. **[CLARIFICATION] Element tag stored in SolutionInterpolation_t is normalised.** Records that the element type in the node payload is the basic (linear) tag of the family — `TETRA_10` and `TETRA_35` are both stored as `TETRA_4` — which is what makes the uniqueness rule well defined, and restates the lookup accordingly (Section 3.4.4). *Rationale:* v2 and earlier v3 text described an exact-tag match followed by a fallback, implying that high-order tags may be found on disk; they cannot, so the fallback is the operative path for any high-order query. Resolution behaviour is unchanged.

New normative sections (no equivalent in v2):

- **[NEW] Mid-Level Library API** — complete C and Fortran 2003 read/write interface for all CPEX-0045 node types, including `LagrangeControlPointDistribution` accessors.
- **[NEW] SIDS File Mapping** — exact on-disk array shapes, data types, and DOF-ordering conventions.
- **[NEW] Implementation Specification** — polynomial order limits, validation requirements, error-code semantics, and default values.
- **[CLARIFICATION] Practical order ceiling of the monomial bases** (Section 3.2.4) — informative note that the monomial mass matrix is Hilbert-like, so conditioning degrades

exponentially with p regardless of normalisation, limiting `ParametricMonomialsPascal` and `CartesianMonomialsPascal` to roughly $p \leq 2-3$ in double precision. Adds no requirement; it bounds the usable range of an already-approved feature and points at the L^2 -orthogonal bases deferred to a future CPEX.

Implementation evidence: The v3 additions are implemented in the CGNS library branch `CPEX45_high_order_wip`: C API (`cgnslib.c/h`), internal readers (`cgns_internals.c`), Fortran bindings (`cgns_f.F90`), and `cgnscheck` validation (including a `-s` “strict CPEX-0045” mode that enforces the order ranges of Section 6). C and Fortran test programs pass. The Cartesian-modal `CharacteristicLength` on-disk node defined in Section 5 is shipped together with the `cg_sol_characteristic_length_{read,write}` helpers in the same branch; the Fortran wrappers follow the standard `_f` suffix convention. Both `CharacteristicLength` encodings, the input validation, and the re-write semantics are covered by the dedicated regression test `src/tests/test_characteristic_length.c`. The bidirectional element-type lookup of Section 3.4.4 is provided by `cg_solution_interpolation_find`.

Deferred to a future CPEX: explicit symbolic representation of individual interpolation functions (noted as a potential extension in v2), which is also the mechanism by which L^2 -orthogonal modal bases (tensor-product Legendre, Dubiner/Koornwinder on simplices) would be described — these are what high-order modal methods require beyond the practical order ceiling of the monomial bases specified here (Section 3.2.4); and space-time mesh geometry (e.g. ALE), which would add a temporal-order attribute to `ElementInterpolation_t` or introduce a dedicated space-time mesh-interpolation node.

Motion. The CGNS Steering Committee moves to adopt **CPEX-0045 version 3.0** (August 1, 2026) as an amendment to the previously approved v2, incorporating: (1) the `LagrangeControlPointDistribution_t` enum and its associated child node; (2) `GridLocation = InterpolationPoints` for all high-order solutions, both uniform-order (whole-zone) and variable-order (`PointRange/PointList` subset); (3) the `InterpolationOrders` child of `FlowSolution_t` as the normative on-disk encoding; (4) mandatory coordinate normalisation for Cartesian modal interpolation, with a normative `CharacteristicLength` child of `FlowSolution_t` supporting both an isotropic (rank-1) and a per-axis (rank-2) encoding, recorded by the writer rather than recomputed by the reader; (5) the Mid-Level Library API, SIDS File Mapping, and Implementation Specification sections as normative components of the standard; and (6) the clarifications listed above, namely that `CharacteristicLength` is excluded from the `FlowSolution_t` field list and that the element tag stored in `SolutionInterpolation_t` is the basic (linear) tag.

Moved by: _____

Seconded by: _____

Vote: For:_____ Against:_____ Abstain:_____ Result:_____

Contents

1	Motivation and Scope	7
2	Rationale	7
3	Extension of the SIDS	7
3.1	Conventions	7
3.1.1	Interpolation Type Enumeration: <code>InterpolationType_t</code>	8
3.1.2	Lagrange Control Point Distribution	9
3.1.3	New <code>GridLocation_t</code> Value: <code>InterpolationPoints</code>	10
3.2	High-Order Parametric Interpolation	11
3.2.1	Standard Coordinate Systems for Parametric Interpolation	11
3.2.2	Parametric Interpolation by Specification of Lagrange Control Points	11
3.2.3	Parametric Modal Interpolation	14
3.2.4	Practical Order Ceiling of the Monomial Bases	15
3.3	Cartesian Modal Interpolation	16
3.3.1	Computation of the Element Coordinate System	16
3.3.2	Cartesian Modal Interpolants	17
3.4	Overriding the Element Definition and Solution Interpolation	18
3.4.1	Modification of Section 12.6 (<code>Family_t</code>)	19
3.4.2	<code>ElementInterpolation_t</code> — New SIDS Section 12.10	19
3.4.3	Changes in Section 7.3 (<code>Elements_t</code>)	20
3.4.4	<code>SolutionInterpolation_t</code> — New SIDS Section 12.11	20
3.4.5	Addition to Section 7.7 (<code>FlowSolution_t</code>)	21
4	Extensions to the Mid-Level Library	22
4.1	Helper Functions	23
4.2	Reading and Writing Interpolation Characteristics via the Family Interface	23
4.3	Accessing Data in the <code>FlowSolution_t</code> Node	25
4.4	Fortran Bindings	26
5	Extension to the SIDS File Mapping	27
5.1	<code>ElementInterpolation_t</code> : Child Node of <code>Family_t</code>	27
5.2	<code>SolutionInterpolation_t</code> : Child Node of <code>Family_t</code>	28
5.3	<code>FlowSolution_t</code> : Added Child Node	29
6	Implementation Specification	29
6.1	Polynomial Order Limits	29
6.2	On-Disk Layout of Control Points and Coefficients	30
6.3	Validation Requirements (Reader)	31
6.4	Cross-Validation Rules (Writer)	32
6.5	Error-Code Semantics	32
6.6	Default Values	32
6.7	Coordinate-System Conventions	33
6.8	Scope of <code>InterpolationPoints</code> Grid Location	33

7	Variable-Order Configurations and IsoParametric Solutions	34
7.1	Variable Order Per Cell (<code>Disjoint PointRange</code>)	34
7.2	Variable Order Per Cell (<code>PointList</code>)	34
7.3	Variable Order Per Field Variable	34
7.4	IsoParametric for <code>SolutionInterpolation_t</code>	35
8	Compliance Tests	35

1 Motivation and Scope

The aim is to cater for the large diversity of high-order methods by including a specification of the interpolation spaces for solutions in the file, fixing the bare essentials: the coordinate system. The choice of Lagrangian interpolation spaces is applicable to the mesh as well as the solution description. This generalisation will:

- allow accurate representation of data in a range of popular interpolation spaces which could otherwise only be approximated;
- allow use of the native storage of the application for the Lagrangian description, leading furthermore to
 - a simpler, more robust, and generic implementation of drivers;
 - more efficient I/O by straightforwardly dumping data blocks;
 - avoiding loss of precision for very specific point distributions;
- avoid the need to redefine the format for each new interpolation type/order;
- cater for space-time methods as well as for ALE computations by including time in the functional expressions.

Currently it is assumed that high-order intra-element interpolation is *restricted to unstructured mesh computations*, as structured meshes do not offer the possibility to individually list elements per order and moreover do not support curved elements. Both modal and nodal interpolations are supported.

There is clear potential for future extension via an explicit specification of mathematical expressions of the individual interpolation functions. This will increase the flexibility for modal interpolation approaches.

2 Rationale

This proposal lifts the limitation of fixing the interpolation functions implicitly by imposing the position of Lagrange control points as proposed for geometric interpolation/curved elements in CPEX 0036 and CPEX 0038. Instead, the description of the functional space and its interpolant is integrated as metadata in the CGNS file in order to allow high flexibility as to the choice of interpolants or even coordinates, an automatic procedure to allow for very high interpolation order, and time-dependent interpolants.

3 Extension of the SIDS

3.1 Conventions

“**Order**” denotes **polynomial degree**. Throughout CPEX-0045, and in v2 before it, *order* means the polynomial *degree* of the expansion — not the number of terms, and not the degree plus one. **SpatialOrder** = p means the interpolation space contains the polynomials of degree $\leq p$: $p = 1$ is linear, $p = 2$ is quadratic. Likewise **TemporalOrder** = $q = 0$ denotes data that is constant in time,

not linear in time. This differs from the common convention under which a degree- p approximation is called order $p + 1$, so a “second-order” method in that sense corresponds to $p = 1$ here.

The v2 usage is retained because `InterpolationOrders`, `SpatialOrder` and `TemporalOrder` are established on-disk node and API argument names, and because v2’s own text is already explicit that what is recorded is a degree — it describes the modal choices as “the maximum *degree* of a Pascal polynomial space” and defines $q = 0$ as the purely spatial case. To keep the two readings from being confused, this document writes “degree” explicitly wherever the numeric value of p or q appears in prose.

Modifications to the SIDS:

- addition of a new section 3.4: *High-order interpolation*.

In the remainder of this section we introduce the different paragraphs (with numbering to be adapted) that should be added to the new section.

“The CGNS standard allows the user to specify their own interpolation approach for both elements and solution. The basic principles are:

1. The element coordinate system per element type is a fixed convention.
2. For the solution interpolation per each element type and interpolation order, a separate interpolation block is added which provides one of three choices:
 - the set of control points for Lagrange interpolants;
 - the maximum degree of a Pascal polynomial space defined in parametric coordinates;
 - the maximum degree of a Pascal polynomial space defined in Cartesian (non-parametric) coordinates.

The first two cover parametric interpolation, whereas the last covers modal/Cartesian interpolation.

3. The mesh is always defined using interpolations in parametric space, by specifying the set of control points. The standard will only allow element types defined up to now, in order not to modify the element connectivity description.
4. The interpolation is *not* supposed to be the same for the geometry as for the solution, unless explicitly specified that way.
5. In addition to the spatial coordinates, time can also be used as an independent variable.

The following sections describe how interpolations are specified.”

3.1.1 Interpolation Type Enumeration: `InterpolationType_t`

“`InterpolationType_t` specifies how the high-order interpolants for the solution are defined. `InterpolationType_t` can take four values:

- **ParametricLagrange** corresponds to a Lagrange interpolation based upon a set of specified control point coordinates in parametric space for a standard interpolation space.

- **ParametricMonomialsPascal** corresponds to modal interpolation functions in parametric coordinates, based on monomials according to the classical Pascal sets.
- **CartesianMonomialsPascal** corresponds to modal interpolation functions in a Cartesian coordinate system, centred on the element and parallel to the main axes. The interpolation functions are based on monomials according to the classical Pascal sets.
- **IsoParametric** corresponds to using the same interpolation functions for the solution as for the mesh.”

3.1.2 Lagrange Control Point Distribution

The parametric coordinates of the control points are stored explicitly in **LagrangeControlPoints** (Section 6.2). Together with the reference domain, fixed by the element type, and the polynomial space, fixed by the element type and order, those coordinates determine the nodal basis uniquely: the functions λ_i satisfying $\lambda_i(\mathbf{u}_j) = \delta_{ij}$ are the solution of a standard multivariate polynomial interpolation problem — obtainable, for instance, by inverting the generalised Vandermonde matrix — and exist and are unique whenever the stored point set is unisolvent for that space. A reader therefore never has to *assume* a distribution, and recording it is not required in order to interpret the file correctly.

Recording it is nonetheless recommended, for three reasons:

- **Conditioning.** A reader that knows the points are Gauss-Lobatto-Legendre (GLL), Gauss-Legendre or Warp&Blend can build the basis from the closed-form barycentric weights of that family, which stays stable at high order, rather than inverting a Vandermonde matrix whose conditioning degrades rapidly as p grows.
- **Precision.** The stored coordinates are R8-rounded, whereas the *interior* nodes of the **GaussLobattoLegendre**, **GaussLegendre** and **WarpAndBlend** families are irrational at all but the lowest orders. Naming the family lets a reader regenerate the nodes at its own working precision instead of inheriting the writer’s rounding.
- **Validation.** **cgnscheck** can compare the stored coordinates against the named family and so flag writer errors — a permuted point list, or a $[0, 1]$ -versus- $[-1, 1]$ domain convention — that are otherwise invisible.

Recommended attribute: **LagrangeControlPointDistribution.** A **DataArray_t** of **Character** type, child of **ElementInterpolation_t** or **SolutionInterpolation_t** when **InterpolationType** is **ParametricLagrange**. Allowed values:

- **GaussLobattoLegendre** (recommended default): 1D points are roots of $(1 - u^2) P'_p(u) = 0$, where P_p is the Legendre polynomial of degree p ; includes endpoints $u_0 = -1$, $u_p = +1$. Optimal conditioning and quadrature accuracy.
- **Equidistant:** $u_i = -1 + 2i/p$. Simple uniform spacing; discouraged for $p > 7$ due to poor conditioning of the Vandermonde matrix.
- **GaussLegendre:** 1D points are roots of $P_p(u) = 0$; does *not* include the endpoints. Rarely used for element interpolation.

- **WarpAndBlend** (simplices only): specialised distribution for triangles and tetrahedra ensuring well-conditioned Vandermonde matrices [4].

Multi-dimensional construction. For tensor-product elements (QUAD, HEXA), the chosen 1D distribution is applied independently in each parametric direction. For simplex elements (TRI, TETRA), **WarpAndBlend** or Fekete distributions are required for $p > 2$; equidistant points may be used at $p = 2$ but become unstable at higher orders.

Precedence: the coordinates are authoritative. **LagrangeControlPoints** and **LagrangeControlPointDistribution** are deliberately redundant, and the redundancy is resolved in favour of the data: the stored coordinates are normative and the distribution name is advisory metadata. A reader *must* construct a basis that interpolates at the coordinates actually present in **LagrangeControlPoints**. It *may* use the named family to obtain those nodes by the better-conditioned or higher-precision route described above, but only after confirming that the family’s nodes agree with the stored coordinates to within its chosen tolerance; if they do not, it *must* fall back to the stored coordinates and *should* warn. It must never silently substitute the named family’s nodes for the stored ones.

This mirrors the rule adopted for **CharacteristicLength** (Section 3.3), where writers record the factors they actually used and readers must use the recorded values rather than recomputing a formula. In both cases the recorded data wins over anything derivable, so that reconstruction never depends on writer and reader agreeing on a computation.

Correspondingly, a writer *must not* record a distribution name that disagrees with the coordinates it writes. **cgnscheck** compares the two whenever both are present and reports a mismatch: an error under **cgnscheck -s**, a warning otherwise.

Absent attribute. If **LagrangeControlPointDistribution** is absent, a reader *must* construct the basis from the stored **LagrangeControlPoints** coordinates directly. It must *not* substitute an assumed distribution, and no warning is warranted, because nothing is ambiguous. An implementation *may* try to recognise the stored points as one of the named families in order to take the better-conditioned path described above, but must fall back to the general construction when no family matches within its chosen tolerance.

3.1.3 New GridLocation_t Value: InterpolationPoints

CPEX-0045 adds a new **GridLocation_t** enumeration value:

- **InterpolationPoints** (enum value 9): the **FlowSolution_t** data is stored at the control points (Lagrange) or as coefficients (modal) defined by the associated **SolutionInterpolation_t** node in the family. The **DataArray_t** length for each field is $\sum_e N_{\text{DOFs}}(e)$, summed over the elements covered by the block (all zone elements when no **PointRange/PointList** is present, or the listed elements otherwise). $N_{\text{DOFs}}(e)$ depends on the interpolation type and orders associated with element e , so for a **MIXED** zone or a block whose listed elements are heterogeneous the per-element DOF count is not constant. DOFs of the first element are stored contiguously, followed by all DOFs of the second element, and so on.

Interpretation of the name. The enumerator is named for the Lagrange case but applies unchanged to the modal types, where the stored values are expansion coefficients not associated with any point. Read `InterpolationPoints` as labelling *the per-element block of interpolation degrees of freedom* — whatever those are for the declared interpolation type — not as an assertion that the data lie at geometric locations. The length rule $\sum_e N_{\text{DOFs}}(e)$ and the element-index semantics below apply identically to both; only the meaning of an individual entry differs.

Mandatory for high-order solutions. For all high-order solutions, `GridLocation = InterpolationPoints` is required, whether the block covers the entire zone (uniform order) or a subset of cells specified by `PointRange` or `PointList` (variable order). Reusing `GridLocation = CellCenter` for high-order data is incorrect: `CellCenter` mandates an array of length equal to the number of cells in the location domain, which contradicts the high-order field-array length $\sum_e N_{\text{DOFs}}(e)$ and will cause failures in standard CGNS readers and `cgnscheck`.

Index semantics under `InterpolationPoints`. When a `FlowSolution_t` carrying `GridLocation = InterpolationPoints` also carries a `PointRange` or `PointList`, the listed indices are interpreted as *element indices* (1-based), not point indices. The expected length of each field `DataArray_t` child is the sum of N_{DOFs} over the listed elements (Section 6.2). In the absence of a point set, the block applies to every element of the zone in element order.

3.2 High-Order Parametric Interpolation

Scope: The parametric interpolation conventions can be used for specifying both curved elements (thereby superseding the standard conventions) and the solution.

3.2.1 Standard Coordinate Systems for Parametric Interpolation

The parametric coordinate system is defined per element following Figure 1.


For space-time computations the coordinate system is extended with one dimension, by the tensor product of the spatial coordinate system with the parameter interval $[-1, 1]$ in parametric time τ . The latter corresponds to the physical time slab $[T_n, T_{n+1}]$.

3.2.2 Parametric Interpolation by Specification of Lagrange Control Points


Scope: Both solution and element shape can be specified using this formulation. For elements, this convention allows redefinition of the position and order of the control points with respect to the standard definitions for each `ElementType_t` described in Section 3.2.1. The standard definition continues to be used when no interpolation is explicitly introduced.

Lagrange interpolants require, next to the specification of control point locations, also the specification of the standard function space $\mathcal{V}$. Its cardinality N defines the number of control points that need to be specified.


We denote the Lagrange interpolant corresponding to control point $\mathbf{u}_i$ as λ_i^1 . Furthermore, we need an arbitrary set of base functions for $\mathcal{V}$, such that $\mathcal{V} = \text{span}(\psi_j, j = 0 \dots N - 1)$. The




(a) TRI




(b) QUAD




(c) TETRA



(d) HEXA



(e) PENTA



(f) PYRA

Figure 1: Parametric coordinate systems per element type.

Lagrange interpolants are then found as a linear combination of the base functions

$$\lambda_i(\mathbf{u}) = \sum_j V_{ij}^{-1} \psi_j(\mathbf{u}) \quad (1)$$

using the inverse of the Vandermonde matrix V associated to the control points $\mathbf{u}_i$ and the basis ψ ,

$$V_{ij} = \psi_j(\mathbf{u}_i). \quad (2)$$

The (spatial) parametric function spaces $\mathcal{V}_p(u, v, w)$ for each element type and order p , which support the Lagrange-type interpolation, are listed in Table 1. Next to the standard “complete” element, also incomplete, or so-called serendipity elements are supported in higher dimensions. A first type of serendipity element only specifies control points on the edges. In three-dimensional elements of sufficient order, a second serendipity interpolation can be defined which only excludes control points internal to the element.

Element		Function space $\mathcal{V}_p(u, v, w)$		
Type	Base type	Complete	Edge Serendipity	Face Serendipity
Line	BAR_2	$\mathcal{L}_p(u)$, $N = p + 1$	n/a	n/a
Quad	QUAD_4	$\mathcal{Q}_p^2(u, v)$, $N = (p + 1)^2$	$\mathcal{L}_p(u) \otimes \mathcal{L}_1(v) \oplus \mathcal{L}_p(v) \otimes \mathcal{L}_1(u)$, $N = 4p$	n/a
Hexa	HEXA_8	$\mathcal{Q}_p^3(u, v, w)$, $N = (p + 1)^3$		
Triangle	TRI_3	$\mathcal{P}_p^2(u, v)$, $N = \frac{(p+1)(p+2)}{2}$		n/a
Tetra	TETRA_4	$\mathcal{P}_p^3(u, v, w)$, $N = \frac{(p+1)(p+2)(p+3)}{6}$		
Prism	PENTA_6	$\mathcal{P}_p^2(u, v) \otimes \mathcal{L}_p(w)$, $N = \frac{(p+1)^2(p+2)}{2}$		
Pyramid	PYRA_5	see [3]		

Table 1: Lagrange functional spaces per element and interpolation type. The base type and order p are specified for solution interpolation; the full element type is used to classify element interpolation.

In which we use direct sums $\oplus$ and products $\otimes$ of the following standard spaces of order p :

- The linear space: $\mathcal{L}_p(u) = \text{span}\{u^i, 0 \leq i \leq p\}$
- Tensor product spaces, $N = (p + 1)^d$:

$$\begin{aligned} \mathcal{Q}_p^2(u, v) &= \mathcal{L}_p(u) \otimes \mathcal{L}_p(v) = \text{span}\{u^i v^j, 0 \leq i, j \leq p\} \\ \mathcal{Q}_p^3(u, v, w) &= \mathcal{L}_p(u) \otimes \mathcal{L}_p(v) \otimes \mathcal{L}_p(w) = \text{span}\{u^i v^j w^k, 0 \leq i, j, k \leq p\} \end{aligned}$$

¹Although for 2D and 3D spaces multi-indices are more convenient, we use for simplicity of notation a single compounded index $i = (i, j, k)$. Likewise a compound coordinate $\mathbf{u} = (u, v, w)$ is used.

- Pascal triangle/tetrahedron, $N = \frac{(p+1)\cdots(p+d)}{d!}$:

$$\begin{aligned}\mathcal{P}_p^2(u, v) &= \text{span}\{u^i v^j, 0 \leq i + j \leq p\} \\ \mathcal{P}_p^3(u, v, w) &= \text{span}\{u^i v^j w^k, 0 \leq i + j + k \leq p\}\end{aligned}$$

For space-time function spaces (e.g. ALE meshes), the complete functional space is given by

$$\mathcal{V}_{p,q}(u, v, w, t) = \mathcal{V}_p(u, v, w) \otimes \mathcal{L}_q(t) \quad (3)$$

where p and q are the spatial and temporal orders respectively.

Prism (PENTA_6) ordering convention. The prism space $\mathcal{P}_p^2(u, v) \otimes \mathcal{L}_p(w)$ is a tensor product of a triangle in (u, v) and a line in w . The implementation requires the line direction w to vary *slowest* when enumerating control points: for each line node along w (w -major), enumerate the $(p+1)(p+2)/2$ triangle nodes in (u, v) . This is *not* a simple 3D tensor product like **HEXA**.

Pyramid (PYRA_5) cardinality. Following Bergot–Cohen–Duruffé [3], the cardinality of the pyramid functional space at order p is

$$N_{\text{PYRA}}(p) = \frac{(p+1)(p+2)(2p+3)}{6}, \quad (4)$$

giving 5, 14, 30, 55 for $p = 1, 2, 3, 4$ respectively. These are the values returned by `cg_element_lagrange_interpolation_size` and consume the corresponding `ElementType_t` tags `PYRA_5`, `PYRA_14`, `PYRA_30`, `PYRA_55`. The control-point layout convention is the user's responsibility; the library does not provide a built-in pyramid layout helper.

3.2.3 Parametric Modal Interpolation

Scope: This type of interpolation only applies to solutions.

For solutions, modal interpolation is also allowed. The interpolation functions are based on an ordered set of monomials spanning the Pascal spaces.

The cardinality of the modal basis is the same for every element shape of a given dimension — the binomial coefficient

$$N_{\text{modal}} = \binom{p+d}{d}, \quad d = \text{element dimension}, \quad p = \text{spatial order, i.e. polynomial degree} \quad (5)$$

For space-time interpolation the cardinality is multiplied by $(q+1)$ where q is the temporal order, again a degree. Consequently a **QUAD_9** element used with modal interpolation of degree $p = 2$ (quadratic; `SpatialOrder = 2`, per Section 3.1) stores $\binom{4}{2} = 6$ monomial coefficients (not 9 — 9 is the Lagrange-basis cardinality of the tensor-product space $\mathcal{Q}_p^2$). The helper functions `cg_element_monomial_size` and `cg_solution_monomial_size` return N_{modal} for any supported element type and order combination.

According to the underlying dimension the monomials are given by:

- 1D: the monomials $1, u, u^2, u^3, \dots, u^p$
- 2D — Pascal triangle ordered as:

```
for (int i=0; i<=p; i++)
  for (int j=0; j<=i; j++)
    f[idx++] = u(i-j)vj
```

- 3D — Pascal tetrahedron ordered as:

```
for (int i=0; i<=p; i++)
  for (int j=0; j<=i; j++)
    for (int k=0; k<=i-j; k++)
      f[idx++] = u(i-j-k)vkwj
```

The above covers purely spatial interpolation. In case the function space is defined in space-time (eg. for ALE meshes), the complete functional spaces are given by

- 1D in space plus time: the spatial monomials multiplied by monomials in (parametric) time τ , ordered in the following way

```
for (int h=0; h<=q; h++)
  for (int i=0; i<=p; i++)
    f[idx++] =  $\tau^h u^i$ 
```

- 2D in space plus time: the spatial Pascal triangle multiplied by monomials in (parametric) time τ , ordered in the following way

```
for (int h=0; h<=q; h++)
  for (int i=0; i<=p; i++)
    for (int j=0; j<=i; j++)
      f[idx++] =  $\tau^h u^{(i-j)} v^j$ 
```

- 3D: the spatial Pascal tetrahedron multiplied by monomials in (parametric) time τ , ordered in the following way

```
for (int h=0; h<=q; h++)
  for (int i=0; i<=p; i++)
    for (int j=0; j<=i; j++)
      for (int k=0; k<=i-j; k++)
        f[idx++] =  $\tau^h u^{(i-j-k)} v^k w^j$ 
```

Note that this space-time formulation reverts to purely spatial interpolation — including the ordering — for the (default) temporal order $q = 0$.

3.2.4 Practical Order Ceiling of the Monomial Bases

Both `ParametricMonomialsPascal` and `CartesianMonomialsPascal` use raw monomials. Implementers should be aware that this choice carries an intrinsic conditioning limit that coordinate

normalisation does *not* remove. The mass matrix of the monomial set $\{1, x, x^2, \dots, x^p\}$ on $[-1, 1]$ has entries $M_{mn} \propto \int_{-1}^1 x^{m+n} dx$, which is a Hilbert-like matrix whose condition number grows exponentially with p . Normalisation (Section 3.3) removes the *scale* disparity between directions and prevents overflow, but it cannot remove this growth.

In practice, modal formulations built on monomials — notably the Taylor-basis discontinuous-Galerkin family [6, 7] — are therefore generally applied at low order, typically $p \leq 2$ or $p \leq 3$ in double precision. Methods that need higher order use L^2 -orthogonal bases instead (tensor-product Legendre on quadrilaterals and hexahedra, Dubiner/Koornwinder on simplices), for which the mass matrix is diagonal. This version of the standard has no way to name such a basis; describing one is the province of the explicit symbolic representation of interpolation functions deferred to a future CPEX. Writers requiring high-order modal data should use `ParametricLagrange` with an appropriate control-point distribution (Section 3.1.2) instead.

3.3 Cartesian Modal Interpolation

Scope: Cartesian modal interpolation only applies to solutions.

3.3.1 Computation of the Element Coordinate System

We specify a local Cartesian coordinate system per element based upon the (simplified) barycenter. Say we note the global coordinates $\mathbf{R} = X\mathbf{e}_x + Y\mathbf{e}_y + Z\mathbf{e}_z$. We proceed by first computing the element barycenter $\mathbf{R}^e$ as the arithmetic mean of the locations of the principal vertices, i.e. the nodes corresponding to those of the associated linear element:

$$\mathbf{R}^e = \frac{1}{N} \sum_{i=1}^N \mathbf{R}_i^e \quad (6)$$

The element local coordinates are then defined as $\mathbf{r} = \mathbf{R} - \mathbf{R}^e$; we then use the notation $\mathbf{r} = x\mathbf{e}_x + y\mathbf{e}_y + z\mathbf{e}_z$.

Correspondingly, an element-local origin of the time dimension is defined as the mid-point of the relevant physical time slab $[T_n, T_{n+1}]$, such that $t = T - (T_n + T_{n+1})/2$.

Numerical stability: mandatory normalization. Using element-local but *unnormalized* coordinates in monomials x^p, y^p, z^p causes catastrophic floating-point overflow or underflow when the physical element size differs greatly from 1 (e.g. kilometres or micrometres), and produces a badly conditioned element mass matrix when the element is strongly stretched. Conformant writers *must* therefore non-dimensionalise the element-local coordinates before forming `MonomialCoefficients`. Two normalisations are permitted.

Isotropic. A single scale factor per element,

$$\xi = x/h^e, \quad \eta = y/h^e, \quad \zeta = z/h^e. \quad (7)$$

This is appropriate for near-equilateral elements and for general polygonal or polyhedral cells, for which axis-aligned extents are not meaningful.

Per-axis (recommended for stretched elements). One scale factor per coordinate direction,

$$\xi = x/h_x^e, \quad \eta = y/h_y^e, \quad \zeta = z/h_z^e. \quad (8)$$

This is the directional non-dimensionalisation used by the Taylor-basis discontinuous-Galerkin family [6, 7], in which the natural choice is the extent of the element along each axis, $h_k^e = \max_{i,j} |\mathbf{e}_k \cdot (\mathbf{R}_i^{\text{vertex}} - \mathbf{R}_j^{\text{vertex}})|$. It is what keeps the modal mass matrix well conditioned on high-aspect-ratio cells such as those used in wall-resolved boundary layers: under a single isotropic h^e dominated by the long axis, the short-direction coordinate collapses toward zero and its higher powers lose relative precision, even though no overflow occurs.

The axes remain parallel to the global Cartesian axes in both cases, as required by the definition of `CartesianMonomialsPascal`; only the scaling differs. Rotated or otherwise oriented local frames are *not* permitted, because the modal coefficients of this basis are interpretable as scaled directional derivatives at the element barycentre, and slope limiters in the Taylor-basis family act on them component-wise along the global axes.

Recorded, not recomputed. Writers *must record the scale factors they actually used*, and readers *must use the recorded values* rather than recomputing any geometric formula. The expressions above are recommended defaults, not requirements: a writer whose solver already works in a normalised frame simply records its own factors and dumps its coefficient blocks unchanged, consistent with the goal (Section 1) of using the native storage of the application without loss of precision. Mandating a formula instead of a record would force such a writer to rescale its coefficients on output, and would leave the reconstruction exposed to any disagreement between the writer’s and the reader’s evaluation of that formula.

On-disk encoding. The scale factors are recorded in a `DataArray_t` child of the parent `FlowSolution_t` node, named `CharacteristicLength`, of R8 datatype. *The array rank selects the normalisation:*

- **rank 1**, shape $[|\mathcal{E}|]$ — isotropic; one h^e per element.
- **rank 2**, shape $[\text{PhysDim}, |\mathcal{E}|]$ — per-axis; `PhysDim` factors per element, with `PhysDim` the fast-varying axis so that the factors of one element are contiguous $(h_x^1, h_y^1, h_z^1, h_x^2, h_y^2, h_z^2, \dots)$, matching the `LagrangeControlPoints` layout convention.

$|\mathcal{E}|$ is the number of elements covered by the block: all elements of the zone for a whole-zone block, or the number of elements listed by `PointRange/PointList` for a subset block. Element ordering within the array matches the element ordering of the corresponding field arrays (Section 5). All factors must be strictly positive. The `CharacteristicLength` child is mandatory when the associated `SolutionInterpolation_t` carries `CartesianMonomialsPascal`, and is otherwise omitted.

Implementation status. The node name and both on-disk shapes are normative. The MLL helpers `cg_sol_characteristic_length_read` and `cg_sol_characteristic_length_write` (Section 4.3) carry an explicit `nscale` argument selecting the encoding, and write and read this node directly. `cgnscheck` validates the datatype, the rank, that `nscale` is either 1 or `PhysDim`, and the element count, whenever the node is present; it also reports a missing `CharacteristicLength` on a `CartesianMonomialsPascal` solution (an error under `cgnscheck -s`, a warning otherwise).

3.3.2 Cartesian Modal Interpolants

The interpolation space for Cartesian modal interpolations follows the same approach as for parametric modal interpolation and is based on ordered sets of monomials spanning the standard Pascal spaces. Here these are monomials in the element-local Cartesian coordinates x, y, z and (physical) time t , rather than coordinates in parametric space and time. The ordered sets of monomials for a purely spatial interpolation are thus given by:

- 1D: the monomials $1, x, x^2, x^3, \dots, x^p$
- 2D — Pascal triangle ordered as:

```
for (int i=0; i<=p; i++)
  for (int j=0; j<=i; j++)
    f[idx++] = x(i-j)yj
```

- 3D — Pascal tetrahedron ordered as:

```
for (int i=0; i<=p; i++)
  for (int j=0; j<=i; j++)
    for (int k=0; k<=i-j; k++)
      f[idx++] = x(i-j-k)ykzj
```

Again, a space-time interpolation can be obtained by multiplying with temporal monomials, which yields the following sets of (ordered) monomials

- 1D in space plus time: the spatial monomials multiplied by monomials in (parametric) time, ordered in the following way

```
for (int h=0; h<=q; h++)
  for (int i=0; i<=p; i++)
    f[idx++] = thxi
```

- 2D in space plus time: the spatial Pascal triangle multiplied by monomials in (parametric) time, ordered in the following way

```
for (int h=0; h<=q; h++)
  for (int i=0; i<=p; i++)
    for (int j=0; j<=i; j++)
      f[idx++] = thx(i-j)yj
```

- 3D: the spatial Pascal tetrahedron multiplied by monomials in (parametric) time, ordered in the following way

```
for (int h=0; h<=q; h++)
  for (int i=0; i<=p; i++)
    for (int j=0; j<=i; j++)
      for (int k=0; k<=i-j; k++)
        f[idx++] = thx(i-j-k)ykzj
```

3.4 Overriding the Element Definition and Solution Interpolation

Modifications to the SIDS:

- include list of solution/element interpolants in section 12.6 `Family_t`;
- new section 12.10 `ElementInterpolation_t`;
- generalisation of section 7.3 `Elements_t` and example in 7.4;
- new section 12.11 `SolutionInterpolation_t`;
- generalisation of section 7.7 `FlowSolution_t` and example in 7.8;
- renumber sections 12.10 `UserDefinedData_t` and 12.11 `Gravity_t`.

3.4.1 Modification of Section 12.6 (`Family_t`)

On a case-by-case basis, i.e. per element type and interpolation order, one can provide alternative mesh and solution interpolants in a dedicated family to the zones in question. This is implemented using dedicated lists of `ElementInterpolation_t` and `SolutionInterpolation_t` leaves within `Family_t`.

```
Family_t :=
```

```
List( Descriptor_t Descriptor1 ... DescriptorN ); (o)
FamilyBC_t FamilyBC ; (o)
...
List( ElementInterpolation_t ElemInterp1 ... ElemInterpN ); (o)
List( SolutionInterpolation_t SolInterp1 ... SolInterpN ); (o)
};
```

3.4.2 `ElementInterpolation_t` — New SIDS Section 12.10

The `ElementInterpolation_t` specifies the geometric interpolation of an element by listing an alternative set of Lagrange high-order control points in parametric space following the element conventions for the coordinate system. In the absence of such a block for a given `ElementType_t`, the standard described in Section 3.2.1 is followed. **It is assumed that the first points correspond to the principal vertices of the corresponding linear element, in the same order (cf. Figure 1).**

```
ElementInterpolation_t :=
```

```
ElementType_t Element; (r)
DataArray_t<Float,DataSize[]> LagrangeControlPoints; (o)
DataArray_t<Float,DataSize[]> MonomialCoefficients; (o)
DataArray_t<Character> LagrangeControlPointDistribution; (o)
};
```

`LagrangeControlPointDistribution` may be present when `LagrangeControlPoints` is present (i.e. the interpolation type is `ParametricLagrange`). It is recommended but not required, and carries no information a reader needs in order to interpret the control points; see Section 3.1.2 and the File Mapping in Section 5.

The interpolation type is not stored explicitly as a child node; it is inferred from which optional children are present:

- `LagrangeControlPoints` present $\Rightarrow$ `ParametricLagrange`
- `MonomialCoefficients` present $\Rightarrow$ `ParametricMonomialsPascal`
- neither present $\Rightarrow$ `IsoParametric` (element's own coordinates used)

`CartesianMonomialsPascal` applies to *solution* interpolation only and cannot be attached to an `ElementInterpolation_t` node.

Limitations: The current proposal maintains `ElementType_t` to describe both element type and geometric order, meaning we cannot go beyond 4th-order interpolation. This choice is motivated by maintaining the `ElementConnectivity_t` leaf in its current shape and the fact that currently there is no real need for higher geometric orders.

Uniqueness. At most one `ElementInterpolation_t` per exact `ElementType_t` tag is permitted within a family. Tags differing only in geometric order are considered distinct; for example a single family may carry separate `QUAD_4` and `QUAD_9` `ElementInterpolation` nodes. The constraint is enforced both on write and on read.

Scope — spatial only. `ElementInterpolation_t` describes *purely spatial* mesh interpolation; it carries no `TemporalOrder` attribute. The space-time function-space machinery of Section 3.2.2 therefore applies to `SolutionInterpolation_t` (which does carry a temporal order) but not to mesh geometry under this version of the standard. Space-time mesh geometry (e.g. ALE) is deferred to a future CPEX, which would add a temporal-order attribute to `ElementInterpolation_t` or introduce a dedicated space-time mesh-interpolation node.

3.4.3 Changes in Section 7.3 (`Elements_t`)

In case an alternative location for the element control points is specified, the actual elements are defined as before by listing the indices in the coordinate table, with the notable change that the order will correspond to the control point coordinates specified in the corresponding `ElementInterpolation_t` block. If a specific element type is not found among these leaves, the standard convention is applied.

3.4.4 `SolutionInterpolation_t` — New SIDS Section 12.11

The interpolation functions associated to the interpolation on a given element type and order are stored in a `SolutionInterpolation_t` leaf attached to the corresponding Family.

`SolutionInterpolation_t` :=

```

int[3] InterpolationData; (r)    /* [ElementType, SpatialOrder, TemporalOrder]
*/
InterpolationType_t InterpolationType; (r)
dataArray_t<Float,DataSize[]> LagrangeControlPoints; (o)
dataArray_t<Float,DataSize[]> MonomialCoefficients; (o)
dataArray_t<Character> LagrangeControlPointDistribution; (o)
};

```

Naming note. v2 referred to this required child informally as `InterpolationName` (see the quoted passage below). The normative on-disk name adopted for v3, consistent with the File Mapping in Section 5 and the reference implementation, is `InterpolationType`. There is only one such child; the two names refer to the same node.

The triple `[ElementType, SpatialOrder, TemporalOrder]` is the node’s own I4 payload (length 3, rank 1), as detailed in Section 5. `TemporalOrder` defaults to 0 when omitted by the writer.

Uniqueness. At most one `SolutionInterpolation_t` per (basic-element-type, `SpatialOrder`, `TemporalOrder`) triplet is permitted within a family.

Stored element tag is the basic tag. The element type in the node’s payload is *normalised to the basic (linear) tag of the element family* when the node is written: a writer given `TETRA_10` or `TETRA_35` stores `TETRA_4`. This is what makes the uniqueness rule above well defined — tags differing only in geometric order are not distinguishable after normalisation, so they cannot coexist — and it is what allows one basis description to serve every geometric order of the same family.

Lookup. A query is therefore resolved in two steps: an exact comparison against the stored tag, and, failing that, a second comparison against the basic tag of the queried type. Both must be tried, because the caller may legitimately present either form. A query for `(TETRA_4, 2, 0)` matches on the first step; a query for `(TETRA_10, 2, 0)` or `(TETRA_35, 2, 0)` matches on the second, and resolves to the same node. Note that, because writers normalise, the first step succeeds only for a basic-tag query; the fallback is the operative path for any high-order tag. A triplet matching neither is reported as absent (`CG_NODE_NOT_FOUND`), which is a normal result and not an error. The MLL function `cg_solution_interpolation_find` (Section 4.2) performs both steps, so callers need not reimplement the fallback.

The relevant `SolutionInterpolation_t` block will be found using the pair composed by the (basic) element type and interpolation order. The former corresponds to either the actual element tag or, if the corresponding `SolutionInterpolation_t` block is absent, the type of the corresponding linear element. For instance, the interpolation functions for the 2nd-order solution on a 4th-order tetrahedron will be associated to element tag `TETRA_35` or `TETRA_4`. Finally, if `InterpolationName` is not specified, standard interpolation (i.e. constant per element) applies.

3.4.5 Addition to Section 7.7 (`FlowSolution_t`)

In case variable high-order solutions are stored, a separate solution block per interpolation order in space (and time) should be stored. The interpolation orders attached to the zone are

indicated by the integers `SpatialOrder` and `TemporalOrder`. The location of the solution is then supposed to be `CellCenter`, and in case of a variable-order solution, one needs to use `PointRange` or `PointList` to single out the elements which will use the specified order.

```
FlowSolution_t<int CellDimension, int IndexDimension,
  int VertexSize[IndexDimension], int CellSize[IndexDimension]> :=
```

```
List( Descriptor_t Descriptor1 ... DescriptorN ); (o)
GridLocation_t GridLocation ; (o/d)
IndexArray_t<1, 2, int> InterpolationOrders ; (o/d)
DataArray_t<R8, 1|2, [NumElements] or [PhysDim,NumElements]>
CharacteristicLength ; (o)
Rind_t<IndexDimension> Rind ; (o/d)
IndexRange<IndexDimension> PointRange ; (o)
IndexArray<IndexDimension, ListLength[], int> PointList ; (o)
List( DataArray_t<DataType, IndexDimension, DataSize[]>
DataArray1 ... DataArrayN ); (o)
DataClass_t DataClass ; (o)
DimensionalUnits_t DimensionalUnits ; (o)
List( UserDefinedData_t UserDefinedData1 ... UserDefinedDataN ); (o)
};
```

The `IndexArray_t` carries `[SpatialOrder, TemporalOrder]` (length 2, datatype I4). `CharacteristicLength` is required only for Cartesian modal interpolation; its rank selects the normalisation (rank 1 isotropic, rank 2 per-axis), as defined in Section 3.3.

The default value for `SpatialOrder` depends on the type of interpolation; the default for `TemporalOrder` is 0.

Implementation note 1 — Accessor functions. The `InterpolationOrders` child is read and written through `cg_sol_interpolation_order_{read,write}`, which take `spatialOrder` and `temporalOrder` as separate scalar arguments while persisting the underlying length-2 `IndexArray_t`. See Section 5 for the on-disk encoding.

Implementation note 2 — Grid location convention. All high-order `FlowSolution_t` nodes use `GridLocation = InterpolationPoints` (see Section 3.1.3), regardless of whether they describe the entire zone or a subset of elements. For variable-order configurations, `PointRange` or `PointList` children of the `FlowSolution_t` list *element* indices, and the per-subset orders are recorded in the `InterpolationOrders` child of each `FlowSolution_t` block. The `cgnscheck` “strict CPEX-0045” mode flags `CellCenter` plus `InterpolationOrders` as a SIDS conflict.

4 Extensions to the Mid-Level Library

Function signature	Modes
<code>ierr = cg_element_lagrange_interpolation_size(ElementType_t t, cgsz_t *sz)</code>	r--
<code>ierr = cg_solution_lagrange_interpolation_size(ElementType_t t, int os, int ot, cgsz_t *sz)</code>	r--
<code>ierr = cg_element_monomial_size(ElementType_t t, cgsz_t *sz)</code>	r--
<code>ierr = cg_solution_monomial_size(ElementType_t t, int os, int ot, cgsz_t *sz)</code>	r--
<code>ierr = cg_element_lagrange_interpolation_count(fn, bn, fam, ElementType_t t, int *cnt)</code>	r--
<code>ierr = cg_solution_lagrange_interpolation_count(fn, bn, fam, ElementType_t t, int os, int ot, int *cnt)</code>	r--

Table 2: Helper functions: interpolation space sizes and family query counts.

Parameter	Description
<code>t</code>	Element type
<code>os</code>	Spatial interpolation order
<code>ot</code>	Temporal interpolation order
<code>sz</code>	Output: cardinality of the interpolation space
<code>cnt</code>	Output: number of matching interpolation nodes in the family

Table 3: Input/output parameters for helper functions.

4.1 Helper Functions

The `*_size` functions return the cardinality of the Lagrange or monomial interpolation space for a given element type and orders. The `*_count` functions query a family to count how many `ElementInterpolation_t` or `SolutionInterpolation_t` nodes of a given type (and order) are already present.

4.2 Reading and Writing Interpolation Characteristics via the Family Interface

The family contains the set of interpolation bases:

- for elements as a function of the element type `ElementType_t`;
- for the solution as a function of the combination `ElementType_t` and two interpolation orders `os` and `ot`. This means that no more than one specification can be present for the triplet (t, os, ot) . The element type always refers back to the baseline element, i.e. the solution interpolation basis for the triplets (TETRA_4,4,0) and (TETRA_10,4,0) are the same;
- the number of coordinates of the control points is defined by the dimension associated to the element type.

IsoParametric mesh interpolation. `cg_element_isoparametric_write` is a convenience writer that creates an `ElementInterpolation_t` node with neither `LagrangeControlPoints` nor

Function signature	Modes
<code>cg_nelement_interpolation_read(fn, bn, fam, *ne)</code>	r--
<code>cg_element_interpolation_read(fn, bn, fam, en, name, *et)</code>	r--
<code>cg_element_interpolation_type_read(fn, bn, fam, en, *it)</code>	r--
<code>cg_element_interpolation_points_read(fn, bn, fam, en, *pu, *pv, *pw)</code>	r--
<code>cg_element_interpolation_write(fn, bn, fam, name, et, *en)</code>	-wm
<code>cg_element_interpolation_points_write(fn, bn, fam, en, *pu, *pv, *pw)</code>	-wm
<code>cg_element_isoparametric_write(fn, bn, fam, name, et, *en)</code>	-wm
<code>cg_nsolution_interpolation_read(fn, bn, fam, *ns)</code>	r--
<code>cg_solution_interpolation_read(fn, bn, fam, sn, name, *et, *os, *ot, *it)</code>	r--
<code>cg_solution_interpolation_points_read(fn, bn, fam, sn, *pu, *pv, *pw, *pt)</code>	r--
<code>cg_solution_interpolation_write(fn, bn, fam, name, et, os, ot, it, *sn)</code>	-wm
<code>cg_solution_interpolation_points_write(fn, bn, fam, sn, *pu, *pv, *pw, *pt)</code>	-wm
<code>cg_solution_interpolation_find(fn, bn, fam, et, os, ot, *sn, *it)</code>	r--
<code>cg_element_interpolation_coefficients_read(fn, bn, fam, en, *coeff)</code>	r--
<code>cg_element_interpolation_coefficients_write(fn, bn, fam, en, *coeff)</code>	-wm
<code>cg_solution_interpolation_coefficients_read(fn, bn, fam, sn, *coeff)</code>	r--
<code>cg_solution_interpolation_coefficients_write(fn, bn, fam, sn, *coeff)</code>	-wm
<code>cg_element_interpolation_distribution_read(fn, bn, fam, en, *dist)</code>	r--
<code>cg_element_interpolation_distribution_write(fn, bn, fam, en, dist)</code>	-wm
<code>cg_solution_interpolation_distribution_read(fn, bn, fam, sn, *dist)</code>	r--
<code>cg_solution_interpolation_distribution_write(fn, bn, fam, sn, dist)</code>	-wm

Table 4: Family-level interpolation API functions.

Param	Description	I/O
<code>fn</code>	CGNS file index number	in
<code>bn</code>	Base index number	in
<code>fam</code>	Family index number	in
<code>ne</code>	Number of element interpolation blocks	in/out
<code>en</code>	Element interpolation index number	in
<code>ns</code>	Number of solution interpolation blocks	in/out
<code>sn</code>	Solution interpolation index number; returned as an output by <code>cg_solution_interpolation_find</code>	in/out
<code>os</code>	Spatial interpolation order	in
<code>ot</code>	Temporal interpolation order	in
<code>it</code>	Interpolation type (<code>InterpolationType_t</code>)	in/out
<code>pu</code>	Control points — <i>u</i> -coordinate	out
<code>pv</code>	Control points — <i>v</i> -coordinate (NULL for 1D)	out
<code>pw</code>	Control points — <i>w</i> -coordinate (NULL for 1D/2D)	out
<code>pt</code>	Control points — time coordinate (NULL if $q = 0$)	out
<code>coeff</code>	Monomial coefficients array (size from <code>cg_*_monomial_size</code>)	in/out
<code>dist</code>	Lagrange control point distribution (<code>LagrangeControlPointDistribution_t</code>)	in/out

Table 5: Input/output parameters for family-level API.

MonomialCoefficients. The interpolation type is implicitly `IsoParametric`: the element's own connectivity table coordinates are reused, and no control-point write call follows. The corresponding read function `cg_element_interpolation_points_read` returns `CG_NODE_NOT_FOUND` on such nodes.

Modal coefficient access. `cg_element_interpolation_coefficients_{read,write}` and the matching `cg_solution_interpolation_coefficients_{read,write}` pair operate on the `MonomialCoefficients` child. The expected array length is given by `cg_element_monomial_size` or `cg_solution_monomial_size`. These functions must not be called on a node whose interpolation type is `ParametricLagrange` or `IsoParametric`.

4.3 Accessing Data in the `FlowSolution_t` Node

Function signature	Modes
<code>cg_sol_ptset_info(fn, bn, zn, sn, *ptype, *npnts)</code>	r--
<code>cg_sol_ptset_read(fn, bn, zn, sn, *pnts)</code>	r--
<code>cg_sol_ptset_write(fn, bn, zn, name, loc, ptype, npnts, *pnts, *sn)</code>	-wm
<code>cg_sol_interpolation_order_read(fn, bn, zn, sn, *os, *ot)</code>	r--
<code>cg_sol_interpolation_order_write(fn, bn, zn, sn, os, ot)</code>	-wm
<code>cg_sol_characteristic_length_read(fn, bn, zn, sn, *nscale, *numElements, *h_e)</code>	r--
<code>cg_sol_characteristic_length_write(fn, bn, zn, sn, nscale, numElements, *h_e)</code>	-wm

Table 6: `FlowSolution`-level interpolation order and characteristic-length API.

Param	Description	I/O
<code>fn</code>	CGNS file index number	in
<code>bn</code>	Base index number	in
<code>zn</code>	Zone index number	in
<code>sn</code>	Solution block index number	in
<code>ptype</code>	Point set type (<code>PointRange</code> or <code>PointList</code>)	out
<code>npnts</code>	Number of points in the point set	out
<code>pnts</code>	Point set data	in/out
<code>loc</code>	Grid location (<code>InterpolationPoints</code> for high-order; both uni-form and variable-order)	in
<code>os</code>	Spatial interpolation order	in/out
<code>ot</code>	Temporal interpolation order	in/out
<code>nscale</code>	Number of normalisation factors per element: 1 selects the isotropic encoding, <code>PhysDim</code> the per-axis encoding	in/out
<code>numElements</code>	Number of elements covered by the solution block	in/out
<code>h_e</code>	Normalisation factors, total length <code>nscale×numElements</code> (NULL on read for a shape-only query)	in/out

Table 7: Input/output parameters for `FlowSolution`-level API.

The interpolation order is assigned per solution block within an unstructured zone, whereas the details concerning the interpolation functions are encoded in the family attached to the zone. In

this case, the solution is supposed to be attached to **CellCenter**; the values are the expansion weights in the basis. The specific interpolation basis is defined through the combination of the element type and the interpolation orders.

The cardinality of the interpolation functions is to be determined first by accessing the description of the interpolation.

4.4 Fortran Bindings

All CPEX-0045 functions are exposed to Fortran via Fortran 2003 BIND(C) interfaces in `cgns_f.F90`, with the conventional `_f` suffix (e.g. `cg_solution_interpolation_write_f`). Argument types follow the standard CGNS Fortran conventions: `integer(csize_t)` for index sizes, `integer(c_int)` for indices and enumerations, and `real(c_double)` for control point and coefficient arrays. A representative Fortran usage example is provided in `src/tests/test_high_orderf.F90`.

5 Extension to the SIDS File Mapping

5.1 ElementInterpolation_t: Child Node of Family_t

Family_t	
Children	
...	
name = <user defined> label = ElementInterpolation_t datatype = I4 dimensions = 1 dimension values = 1 data = <element type> cardinality = 0:N	
Children	Comments
name = LagrangeControlPoints type = DataArray_t datatype = R8 IndexDimension = 2 DataSize[] = [Dimension, NumberOfPoints] data = <point locations> cardinality = 0:1	IndexDimension is the rank of the array (2). Dimension is the spatial dimension of the element.
name = MonomialCoefficients datatype = R8 data = <monomial coefficients> cardinality = 0:1	size from cg_element_monomial_size . Present for ParametricMonomialsPascal .
name = LagrangeControlPointDistribution datatype = C1 cardinality = 0:1	Name of the distribution (e.g., GaussLobattoLegendre). Recommended, not required, when interpolation type is ParametricLagrange .

Table 8: Node structure for ElementInterpolation_t.

5.2 SolutionInterpolation_t: Child Node of Family_t

Family_t	
Children	
...	
name = <user defined> type = SolutionInterpolation_t datatype = I4 dimensions = 1 dimension values = 3 data = <element type, spatial order, temporal order> cardinality = 0:N	
<i>The element type is the basic (linear) tag of the element family, normalised on write: TETRA_10 and TETRA_35 are both stored as TETRA_4 (Section 3.4.4).</i>	
Children	Comments
name = InterpolationType datatype = InterpolationType_t data = <choice for interpolation type>	ParametricLagrange, ParametricMonomialsPascal, CartesianMonomialsPascal or IsoParametric
name = LagrangeControlPoints type = DataArray_t datatype = R8 IndexDimension = 2 DataSize[] = [Dimension', NumberOfPoints] data = <point locations> cardinality = 0:1	IndexDimension is the rank of the array (2). Dimension' = element Dimension for purely spatial, or Dimension+1 for space-time (extra row holds parametric time τ). NumberOfPoints = $N_{\text{spatial}} \times (q + 1)$ for space-time.
name = MonomialCoefficients datatype = R8 data = <monomial coefficients> cardinality = 0:1	size from cg_solution_monomial_size. Present for ParametricMonomialsPascal or CartesianMonomialsPascal.
name = LagrangeControlPointDistribution datatype = C1 cardinality = 0:1	Name of the distribution (e.g., GaussLobattoLegendre). Recommended, not required, when interpolation type is ParametricLagrange.

Table 9: Node structure for SolutionInterpolation_t.

5.3 FlowSolution_t: Added Child Node

A single child node is added to `FlowSolution_t`:

FlowSolution_t
Children
...
name = InterpolationOrders type = IndexArray_t datatype = I4 dimensions = 1 dimension values = 2 data = <spatial and temporal interpolation order> cardinality = 0:1
name = CharacteristicLength type = DataArray_t datatype = R8 dimensions = 1 or 2 (rank selects the normalisation) dimension values = [$\mathcal{E}$] isotropic: one factor per element [PhysDim, $\mathcal{E}$] per-axis: PhysDim factors per element, PhysDim fast-varying data = <per-element normalisation factors> cardinality = 0:1
<i>Mandatory when the associated SolutionInterpolation_t has InterpolationType = CartesianMonomialsPascal; element ordering matches the field-array ordering. All factors must be strictly positive. $\mathcal{E}$ is the number of elements covered by the block. Writers record the factors they actually used; readers use the recorded values and do not recompute them (Section 3.3).</i>

Table 10: New InterpolationOrders and CharacteristicLength children of FlowSolution_t.

6 Implementation Specification

This section consolidates the formal constraints, on-disk encodings, error semantics, and default values that a conforming implementation must observe. It applies to both readers and writers.

6.1 Polynomial Order Limits

- **Geometric order** (encoded via `ElementType_t`): bounded by the available high-order tags, currently up to fourth order (e.g. `HEXA_125`, `TETRA_35`).
- **Solution spatial order**: non-negative integer. The library writer (`cg_sol_interpolation_order_write`) accepts any value ≥ 0 . `cgnscheck` warns when the spatial order falls outside the typical range $[0, 100]$; in strict CPEX-0045 mode (`cgnscheck -s`) values outside $[1, 100]$ are reported as errors and `SpatialOrder = 0` is also flagged as an error because no valid Lagrange basis corresponds to it.
- **Solution temporal order**: non-negative integer. `cgnscheck` warns when outside $[0, 10]$; strict mode reports values outside $[0, 10]$ as errors.

- **Constraint:** `TemporalOrder > 0` requires `SpatialOrder ≥ 1` (enforced by `cg_sol_interpolation_order_write`).

6.2 On-Disk Layout of Control Points and Coefficients

LagrangeControlPoints (purely spatial). The data array is stored as a 2D R8 array with shape $[Dim, NPoints]$ in CGNS (Fortran-style) ordering, i.e. Dim is the fast-varying axis and $NPoints$ the slow-varying axis. Equivalently, the byte stream is the concatenation of $NPoints$ tuples, each of length Dim :

$$(u_0, v_0, w_0, u_1, v_1, w_1, \dots, u_{N-1}, v_{N-1}, w_{N-1})$$

LagrangeControlPoints (space-time). For `SolutionInterpolation_t` with `TemporalOrder > 0`, the array shape becomes $[Dim + 1, NPoints_{ST}]$ where $NPoints_{ST} = NPoints_{spatial} \times (q + 1)$ and $q = TemporalOrder$. Within each coordinate row, the points are ordered *time-major*:

$$idx = t \times N_{spatial} + s, \quad t = 0, \dots, q, \quad s = 0, \dots, N_{spatial} - 1$$

That is, all spatial points at time level $t = 0$ are listed first, then all spatial points at $t = 1$, and so on. Because the array has shape $[Dim + 1, NPoints_{ST}]$ in CGNS Fortran-style (column-major) ordering, the $Dim + 1$ dimension is the fast-varying axis. Consequently the byte stream is the concatenation of $NPoints_{ST}$ interleaved tuples, each of length $Dim + 1$:

$$(u_0, v_0, w_0, \tau_0, u_1, v_1, w_1, \tau_1, \dots)$$

The spatial components and the temporal component τ are therefore *interleaved per point*; τ does *not* appear as a contiguous trailing block. Solution-data arrays in `FlowSolution_t` for space-time interpolation follow the same time-major ordering within each element's $N_{DOFs}(e)$ block.

MonomialCoefficients. A 1D R8 array of length $N_{modal} = \binom{p+d}{d}$ for purely spatial interpolation, or $N_{modal} \times (q + 1)$ for space-time, where d is the element dimension, p is the spatial order, and q is the temporal order. The Pascal-traversal order is given by the loops in Section 3.2.3 and is a writer convention; the library does not validate ordering.

InterpolationOrders. A 1D `IndexArray_t` (data type I4) of length 2, child of `FlowSolution_t`, holding `[spatialOrder, temporalOrder]`. Recognized only for `Unstructured` zones.

High-order FlowSolution_t field arrays. When `GridLocation = InterpolationPoints`, each field `DataArray_t` child of `FlowSolution_t` has length

$$L = \sum_{e \in \mathcal{E}} N_{DOFs}(e),$$

where $\mathcal{E}$ is the set of elements covered by the block (all zone elements when no `PointRange/PointList` is present, or the listed elements otherwise) and $N_{DOFs}(e)$ is the per-element DOF count from the matching `SolutionInterpolation_t` for element e . For a homogeneous block this

collapses to $|\mathcal{E}| \cdot N_{\text{DOFs}}$. Within the array, the DOFs of the first listed element are stored contiguously, followed by all DOFs of the second listed element, and so on:

$$\text{array index} = \left(\sum_{k < i} N_{\text{DOFs}}(e_k) \right) + \text{local_dof_index}$$

where e_i is the i -th listed element (0-based) and local_dof_index runs from 0 to $N_{\text{DOFs}}(e_i) - 1$ in the ordering established by the associated `SolutionInterpolation_t` (Lagrange control points listed in `LagrangeControlPoints`, or modal coefficients in the Pascal traversal of Section 3.2.3).

Metadata arrays are not solution fields. `CharacteristicLength` is a `DataArray_t` child of `FlowSolution_t`, but it is interpolation *metadata* and carries its own shape rather than the field length L above. Readers that enumerate the `DataArray_t` children of a `FlowSolution_t` to build the list of solution fields *must exclude it by name*, and must not apply the field size rule to it. A reader that does not exclude it will both report a spurious extra field and reject a conformant file, because the metadata length ($\text{n scale} \times |\mathcal{E}|$) does not equal L . Implementations should note that this failure is invisible on test files that contain no element sections, since field size checking cannot be performed at all in that case.

CharacteristicLength. An R8 `DataArray_t` child of `FlowSolution_t` holding the per-element coordinate normalisation factors used by Cartesian modal interpolation (Section 3.3). The array rank selects the normalisation: rank 1 with shape $|\mathcal{E}|$ is isotropic (one factor per element); rank 2 with shape $[\text{PhysDim}, |\mathcal{E}|]$ is per-axis (PhysDim factors per element, PhysDim fast-varying, so one element's factors are contiguous). All factors must be strictly positive. Mandatory when the associated `SolutionInterpolation_t` uses `CartesianMonomialsPascal`; absent otherwise. Element ordering matches the field-array element ordering above. The values recorded are those the writer actually applied; readers must not recompute them from element geometry.

6.3 Validation Requirements (Reader)

A conforming reader must reject:

- `ElementInterpolation_t` payloads that are not I4 scalar (the `ElementType_t` value);
- `SolutionInterpolation_t` payloads that are not I4, dimension 1, with `dim_vals[0] = 3` (the triple $[\text{ElementType}, \text{SpatialOrder}, \text{TemporalOrder}]$);
- `LagrangeControlPoints` children whose data type is not R8 or whose array dimension is not 2;
- `MonomialCoefficients` children whose data type is not R8 or whose array dimension is not 1;
- `DataArray_t` children of `ElementInterpolation_t` or `SolutionInterpolation_t` whose name is anything other than `LagrangeControlPoints`, `MonomialCoefficients` or `LagrangeControlPointDistribution`;
- `SolutionInterpolation_t` nodes that do not contain exactly one `InterpolationType_t` child named `InterpolationType`;
- Duplicate `ElementInterpolation_t` nodes (same `ElementType_t` tag) or duplicate `SolutionInterpolation_t` nodes (same basic-type, `os`, `ot`) within a family;

- `CharacteristicLength` children whose data type is not `R8`, whose array rank is neither 1 nor 2, whose `nscale` (rank-2 only) is neither 1 nor `PhysDim`, whose element count disagrees with the block, or which contain a non-positive factor.

A conforming reader must also *exclude* `CharacteristicLength` from the solution-field list of its parent `FlowSolution_t`, and must not apply the field size rule of Section 6.2 to it.

A mismatch between `LagrangeControlPointDistribution` and the coordinates in `LagrangeControlPoints` is *not* grounds for rejection. The coordinates are authoritative (Section 3.1.2): a conforming reader must build the basis from them and should warn about the disagreement, rather than failing the file or adopting the named family's nodes.

6.4 Cross-Validation Rules (Writer)

- `cg_solution_interpolation_points_write` must be rejected when the parent `SolutionInterpolation_t` has `InterpolationType` of `ParametricMonomialsPascal` or `CartesianMonomialsPascal` (modal nodes do not carry Lagrange points; coefficients must be written instead).
- `cg_element_isoparametric_write` must not be followed by `cg_element_interpolation_points_write` or `cg_element_interpolation_coefficients_write` on the same node.
- Re-writing an existing `LagrangeControlPoints` or `MonomialCoefficients` array requires the file to be opened in `CG_MODE_MODIFY`; in `CG_MODE_WRITE` the second write must be rejected.
- A writer must not record a `LagrangeControlPointDistribution` whose family nodes disagree with the `LagrangeControlPoints` coordinates written on the same node. `cgnscheck` compares the two whenever both are present and reports a mismatch — an error under `cgnscheck -s`, a warning otherwise. Because the coordinates are authoritative (Section 3.1.2), the remedy is to correct or omit the name, never to alter the coordinates to match it.
- `cg_sol_interpolation_order_write` requires `GridLocation = InterpolationPoints` on the parent `FlowSolution_t`. For backward compatibility with files written under earlier drafts, the library *also* accepts `GridLocation = CellCenter` provided a `PointRange` or `PointList` is present, but `cgnscheck -s` flags this combination as non-conformant and conformant writers must use `InterpolationPoints` exclusively.

6.5 Error-Code Semantics

Detecting interpolation type at runtime. For `ElementInterpolation_t`, call `cg_element_interpolation_type_read`; the returned value is one of `ParametricLagrange`, `ParametricMonomialsPascal`, or `IsoParametric`. `CartesianMonomialsPascal` is never returned for mesh interpolation. For `SolutionInterpolation_t`, the type is the `InterpolationType` child read by `cg_solution_interpolation_read`.

6.6 Default Values

- `SolutionInterpolation_t`: `SpatialOrder` is required (no default); `TemporalOrder` defaults to 0.

Code	Meaning
CG_OK	Operation succeeded.
CG_NODE_NOT_FOUND	Returned by <code>cg_element_interpolation_points_read</code> and <code>cg_solution_interpolation_points_read</code> when the corresponding node has no <code>LagrangeControlPoints</code> child (i.e. the interpolation type is <code>IsoParametric</code> or <code>modal</code>); also returned by <code>cg_sol_interpolation_order_read</code> when the <code>FlowSolution_t</code> has no <code>InterpolationOrders</code> child. Callers must treat this as a normal control-flow signal, not an error.
CG_ERROR	Validation failure or I/O error. Use <code>cg_get_error()</code> to retrieve the diagnostic message string.

Table 11: Error-code semantics for CPEX-0045 functions.

- `FlowSolution_t.InterpolationOrders`: when the child is absent `cg_sol_interpolation_order_read` returns `CG_NODE_NOT_FOUND` with output values (0, 0). Such a `FlowSolution_t` is treated as a standard (non-high-order) solution.

6.7 Coordinate-System Conventions

All supported elements use the *bi-unit* reference domain, i.e. each parametric coordinate ranges over $[-1, 1]$ rather than $[0, 1]$. These are the standard reference elements of the high-order literature [5], and they are the domains drawn in Figure 1, which is normative:

- Tensor-product elements (BAR, QUAD, HEXA): the bi-unit cube $[-1, 1]^d$.
- Triangle (TRI): $\{(u, v) : u \geq -1, v \geq -1, u + v \leq 0\}$, with vertices $(-1, -1)$, $(1, -1)$, $(-1, 1)$. The hypotenuse is the line $u + v = 0$.
- Tetrahedron (TETRA): $\{(u, v, w) : u, v, w \geq -1, u + v + w \leq -1\}$, with vertices $(-1, -1, -1)$, $(1, -1, -1)$, $(-1, 1, -1)$, $(-1, -1, 1)$. The oblique face is the plane $u + v + w = -1$.
- Prism (PENTA): the TRI cross-section above, in (u, v) , extruded along $w \in [-1, 1]$.
- Pyramid (PYRA): $\{(u, v, w) : -1 \leq w \leq 1, |u| \leq \frac{1-w}{2}, |v| \leq \frac{1-w}{2}\}$ — the square base $[-1, 1]^2$ at $w = -1$ contracting linearly to the apex $(0, 0, 1)$. The functional space on this domain is not a polynomial space and is defined in [3]; its cardinality is given in Section 3.2.2.

This is an interoperability requirement, not a recommendation. Control points are stored as bare parametric coordinates, so a writer and a reader that assume different reference domains will exchange a file that is structurally valid and passes `cgnscheck` while being geometrically wrong. A $[0, 1]$ -based simplex convention in particular is *not* conformant: its coordinates are indistinguishable from valid bi-unit coordinates, so the error cannot be detected after the fact. The library does *not* enforce these ranges on user-supplied control points — any `double` value is accepted — which is precisely why conformance is the writer’s responsibility.

6.8 Scope of InterpolationPoints Grid Location

`GridLocation_t = InterpolationPoints` is valid only on `FlowSolution_t` nodes. `ArbitraryGridMotion_t` and `DiscreteData_t` nodes must not use this location.

7 Variable-Order Configurations and IsoParametric Solutions

This section describes the three usage patterns that arise when different cells or different field variables require different polynomial orders, and the use of `IsoParametric` for solution interpolation.

7.1 Variable Order Per Cell (Disjoint PointRange)

When different subsets of cells use different polynomial orders, create one `FlowSolution_t` block per subset using `cg_sol_ptset_write` with non-overlapping `PointRange` intervals, each annotated with `cg_sol_interpolation_order_write`:

```
/* Cells 1..4 at order 2, cells 5..8 at order 3 */
cgsizet range_p2[2] = {1, 4};
cg_sol_ptset_write(fn,B,Z,"FS_p2",InterpolationPoints,PointRange,2,
    range_p2,&S1);
cg_sol_interpolation_order_write(fn,B,Z,S1, 2,0);

cgsizet range_p3[2] = {5, 8};
cg_sol_ptset_write(fn,B,Z,"FS_p3",InterpolationPoints,PointRange,2,
    range_p3,&S2);
cg_sol_interpolation_order_write(fn,B,Z,S2, 3,0);
```

7.2 Variable Order Per Cell (PointList)

Non-contiguous subsets require `PointList` instead of `PointRange`:

```
cgsizet odd[4]   = {1,3,5,7};
cgsizet even[4]  = {2,4,6,8};
cg_sol_ptset_write(fn,B,Z,"FS_Odd", InterpolationPoints,PointList,4,odd,
    &S1);
cg_sol_interpolation_order_write(fn,B,Z,S1, 2,0);
cg_sol_ptset_write(fn,B,Z,"FS_Even",InterpolationPoints,PointList,4,
    even,&S2);
cg_sol_interpolation_order_write(fn,B,Z,S2, 3,0);
```

7.3 Variable Order Per Field Variable

When different physical fields use different polynomial orders over the *same* set of cells, create one `FlowSolution_t` per field variable with the same `PointRange` but distinct interpolation orders:

```
cgsizet all[2] = {1, N_ELEM};
/* Density at order 2 */
cg_sol_ptset_write(fn,B,Z,"FS_Density",InterpolationPoints,PointRange,
    2,all,&S1);
cg_sol_interpolation_order_write(fn,B,Z,S1, 2,0);
cg_field_write(fn,B,Z,S1,RealDouble,"Density",density,&fld);

/* VelocityX at order 3 */
```

```
cg_sol_ptset_write(fn,B,Z,"FS_VelocityX",InterpolationPoints,PointRange
,2,all,&S2);
cg_sol_interpolation_order_write(fn,B,Z,S2, 3,0);
cg_field_write(fn,B,Z,S2,RealDouble,"VelocityX",velocity,&fld);
```

7.4 IsoParametric for SolutionInterpolation_t

When `InterpolationType` is set to `IsoParametric`, the solution uses the same interpolation basis as the mesh (i.e. the `ElementInterpolation_t` defined for the same element type in the family). No `LagrangeControlPoints` or `MonomialCoefficients` child is stored. The corresponding `ElementInterpolation_t` must exist in the family.

```
int sn;
cg_solution_interpolation_write(fn,bn,fam,"QUAD4_IsoParam",
    QUAD_4, /*os=*/1, /*ot=*/0, IsoParametric, &sn);
/* No cg_solution_interpolation_points_write call needed */
```

Reading back verifies that `cg_solution_interpolation_points_read` returns `CG_NODE_NOT_FOUND` for an `IsoParametric` node.

8 Compliance Tests

A conforming implementation must pass the following objective tests. These are independent of any particular library and can be applied directly to a written file.

Partition of unity (Lagrange). For `ParametricLagrange` interpolation of any supported element type and order p ,

$$\sum_{i=1}^N \lambda_i(\mathbf{u}) = 1 \quad \text{for any } \mathbf{u} \text{ in the parametric domain,} \quad (9)$$

within tolerance $\epsilon = 10^{-12}$. Reference test points, each strictly interior to its reference domain (Section 6.7): **QUAD** $p = 2$ at $(u, v) = (0.3, -0.7)$; **HEXA** $p = 2$ at $(0.5, -0.5, 0.2)$; **TRI** $p = 2$ at $(-0.5, -0.3)$, for which $u + v = -0.8 < 0$; **TETRA** $p = 2$ at $(-0.6, -0.5, -0.4)$, for which $u + v + w = -1.5 < -1$.

Kronecker delta (Lagrange). At each control point $\mathbf{u}_j$,

$$\lambda_i(\mathbf{u}_j) = \delta_{ij} \quad (10)$$

within tolerance 10^{-12} for all $i, j \in \{1, \dots, N\}$.

Modal coefficient ordering. Represent a known polynomial in the chosen modal basis and verify that the resulting coefficient vector matches the expected ordering of Section 3.2.3 (parametric Pascal traversal) or Section 3.3 (Cartesian column-major). For example, the polynomial $f(u, v) = u^2 + v$ over a **QUAD** $p = 2$ Cartesian modal basis must produce a coefficient vector whose only non-zero

entries are $c[u^2v^0] = 1$ and $c[u^0v^1] = 1$; if the indices of those non-zero entries differ from the convention, the implementation is non-compliant.

Geometry reconstruction. For a curved `QUAD_9` element with the centre node displaced from its parametric midpoint, evaluating the interpolated geometry at $(u, v) = (0, 0)$ must reproduce the displaced centre coordinates within tolerance 10^{-10} .

Numerical stability (Cartesian modal). For a `QUAD` element whose physical extents span large or small scales (e.g. 10^6 or 10^{-6} in either coordinate), modal-coefficient writes followed by reads must round-trip the data without overflow or precision loss when the writer applies the normalization of Section 3.3. Failure of this test indicates missing or incorrect normalization at the writer. The reader must additionally verify that the normalisation factors recovered from the `CharacteristicLength` node correctly reconstruct the unnormalised coordinate space without precision loss, for both the isotropic (rank-1) and per-axis (rank-2) encodings.

Anisotropic conditioning (Cartesian modal, per-axis). For a strongly stretched element — e.g. a `QUAD` or `HEXA` of aspect ratio 10^4 , representative of a wall-resolved boundary-layer cell — a per-axis `CharacteristicLength` must round-trip exactly and must place every normalised coordinate in $O([-1, 1])$. An implementation that emits a single isotropic factor for such an element is conformant only if it accepts the resulting loss of relative precision in the short direction; per-axis normalisation is the recommended encoding in this case.

Compliance checklist.

- Partition-of-unity test for every supported element type and order.
- Kronecker delta property at every control point.
- Modal-coefficient ordering matches Pascal traversal (parametric) or column-major tensor product (Cartesian).
- Geometry reconstruction within tolerance 10^{-10} .
- Numerical stability for extreme coordinate scales.
- Time-major ordering for space-time data layout.
- Principal vertices identified per the linear sub-element rule (Cartesian modal only).
- Coordinate normalization applied for Cartesian modal interpolation, with the factors actually used recorded in `CharacteristicLength`.
- Both `CharacteristicLength` encodings handled on read: rank-1 isotropic and rank-2 per-axis.
- Lagrange basis constructed from the stored `LagrangeControlPoints` coordinates whether or not a `LagrangeControlPointDistribution` attribute is present, and no distribution assumed when it is absent.

References

-
- [1] CPEX 0038: “Quartic Elements for High Order”.
 - [2] CPEX 0036: “Cubic Elements for High Order”.
 - [3] M. Bergot, G. Cohen, and M. Duruflé, “Higher-order Finite Elements for Hybrid Meshes Using New Nodal Pyramidal Elements”, *Journal of Scientific Computing*, vol. 42, pp. 345–381, 2010. [doi:10.1007/s10915-009-9334-9](https://doi.org/10.1007/s10915-009-9334-9)
 - [4] T. Warburton, “An explicit construction of interpolation nodes on the simplex”, *Journal of Engineering Mathematics*, vol. 56, no. 3, pp. 247–262, 2006. [doi:10.1007/s10665-006-9086-6](https://doi.org/10.1007/s10665-006-9086-6)
 - [5] J. S. Hesthaven and T. Warburton, *Nodal Discontinuous Galerkin Methods: Algorithms, Analysis, and Applications*, Springer Texts in Applied Mathematics, vol. 54, Springer, 2008. [doi:10.1007/978-0-387-72067-8](https://doi.org/10.1007/978-0-387-72067-8)
 - [6] H. Luo, J. D. Baum, and R. Löhner, “A discontinuous Galerkin method based on a Taylor basis for the compressible flows on arbitrary grids”, *Journal of Computational Physics*, vol. 227, no. 20, pp. 8875–8893, 2008. [doi:10.1016/j.jcp.2008.06.035](https://doi.org/10.1016/j.jcp.2008.06.035)
 - [7] D. Kuzmin, “Entropy stabilization and property-preserving limiters for discontinuous Galerkin discretizations of nonlinear hyperbolic equations”, *Journal of Numerical Mathematics*, vol. 29, no. 4, pp. 307–322, 2021. [arXiv:2004.03521](https://arxiv.org/abs/2004.03521)
 - [8] S. Massago, *mathex* library, part of the Small Scientific Library (SSCILIB). <http://sscilib.sourceforge.net>